

MySqlWorkbench Tutorial:

Creating Related Database Tables

(Primary Keys, Foreign Keys, Joining Data)

Contents

1. Overview	2
2. Before You Start	2
3. Connect to MySQL using MySqlWorkbench	2
4. Create Tables “web_user” and “user_role”	2
5. Add Foreign Key from Table “web_user” to Table “user_role”	4
6. Add Data to Table “user_role” (and Get Primary Key Violation)	5
7. Add Data to Table “web_user” (and Get Foreign Key Violation)	6
8. SELECT Data	7
9. Join Data	9

1. Overview

After reviewing database terms, you will use MySQLWorkbench to create a database with 2 tables (customer and product) and enter test data into both of these tables. You will learn about primary keys (auto-increment and not auto-increment) and primary key constraints. You will also learn how to write and execute simple (single table) SQL select statements.

2. Before You Start

There is a separate document that tells you how to **install MySQLWorkbench** on your PC/MAC. MySQLWorkbench is a GUI (Graphical User Interface) front end to the MySQL database management system (which runs on a temple database server).

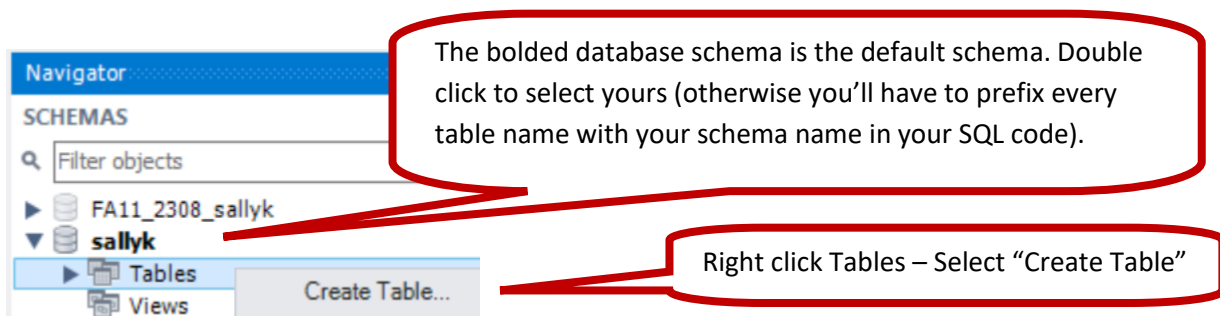
There is a separate document which **reviews database terms and concepts**.

3. Connect to MySQL using MySQLWorkbench

There is a separate document that tells you how to find your database credentials and connect to the database that was created for your use in your class. It's a little involved the first time since you have to find your database password, etc.

4. Create Tables "web_user" and "user_role"

Using MySQLWorkbench, locate then double click on your database schema (left pane under "SCHEMAS"). Your schema will then become bold indicating that this is the "default database". Then open up your schema and right click on "Tables" (under your database schema), then select "create table".



From the “New Table” window, enter the new table name (“web_user”), check that “InnoDB” is selected for database engine “InnoDB” (so foreign keys will work), and open up the two chevrons so that there is space to enter in the columns for the web_user table.

Mouse over the abbreviations above the check boxes to learn the various attributes of database table columns:

- **PK:** Primary Key (unique identifier of each record within the table)
- **NN:** Not Null (not allowed to be null, means it is a required field for the user to enter)
- **UQ:** field must be unique within the table.
- **AI:** Auto-increment. The database management system supplies the next available number (automatically) – the user does not have to enter a value into this field.

Then enter in the column names, data types, and attributes of the web_user table exactly as shown. You will be getting sample code all semester that expects this table to be just so.

Column Name	Datatype	PK	NN	UQ	B	UN	ZF	AI	G	Default/Expression
web_user_id	INT(11)	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	
user_email	VARCHAR(55)	<input type="checkbox"/>	<input type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	NULL
user_password	VARCHAR(45)	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	NULL
birthday	DATE	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	NULL
membership_fee	DECIMAL(12,2)	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	NULL
user_role_id	INT(11)	<input type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	

Create a second table named “user_role” with columns and attributes (exactly) as shown:

Column Name	Datatype	PK	NN	UQ	B	UN	ZF	AI	G	Default/Expression
user_role_id	INT(11)	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	
user_role_type	VARCHAR(10)	<input type="checkbox"/>	<input checked="" type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	<input type="checkbox"/>	

If you ever need to **modify the design of one of your tables**, right click on the table, then select “**Alter Table**”.

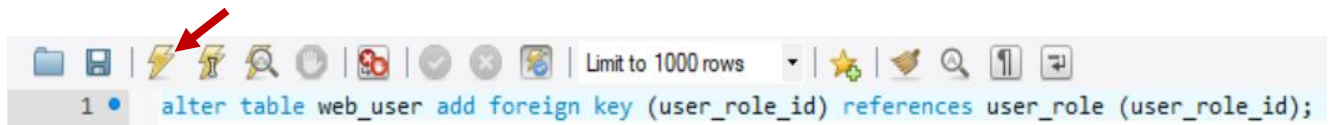
If you are **not seeing a table** that you just created, right click on “Tables” and select “**Refresh All**”.

When you design a database table,

- Typically you will have one column that is the primary key (PK, NN). By naming convention, the name of this field should be tableName_id. A primary key may often be auto-increment (AI) which means that the database management system is “in charge of” allocating a unique number to each record as the records are inserted. In this tutorial, one of the tables (web_user) has an auto-increment primary key, but the other (user_role) does not. This is to show you the difference between the two typical design choices.
- For other columns, specify the desired data type. Here are some of the “normal” types are:
 - VARCHAR: specify the maximum size of a string to be stored.
 - INT: this is a normal integer. Specify the maximum number of digits to be stored.
 - DECIMAL: this is the best data type to store currency. It’s stored like an integer (not real number/floating point number), but with the decimal point moved to the left a certain number of digits. Specify the maximum digits and the number of digits to the right of the decimal point.
 - DATE: this is the data type to use (not DATETIME) if you just want to store day/month/year. DATETIME is more complicated to use – it includes hours/minutes/seconds/milliseconds and it requires more effort to convert and format.
- In our data model, we want each web_user record to indicate which user_role they are. To implement this, we add field user_role_id to table web_user (type matching user_role.user_role_id) and in the next step we will turn that field into a foreign key (like a pointer). It is a database naming convention to name a foreign key field the same as the primary key that it references.

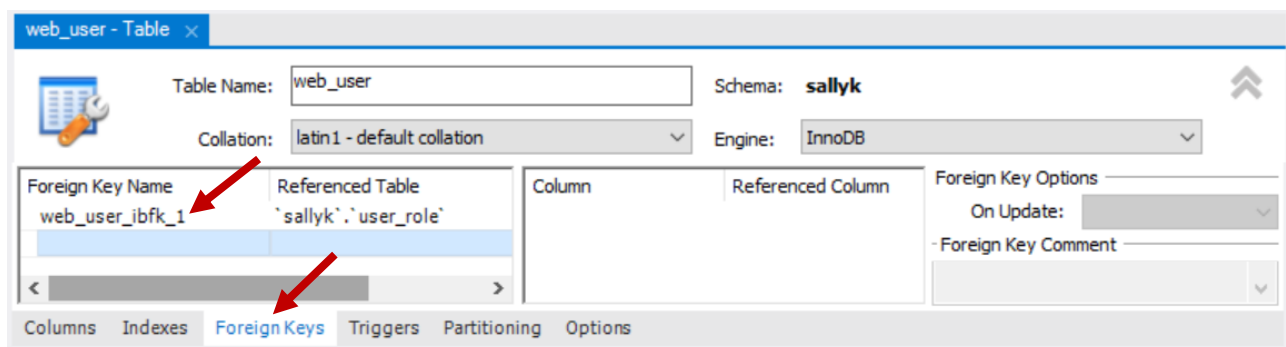
5. Add Foreign Key from Table “web_user” to Table “user_role”

A foreign key is like a pointer. To create a foreign key from web_user.user_role_id to user_role.user_role_id, type the following SQL into the Query area then click on the execute icon (lightning bolt).



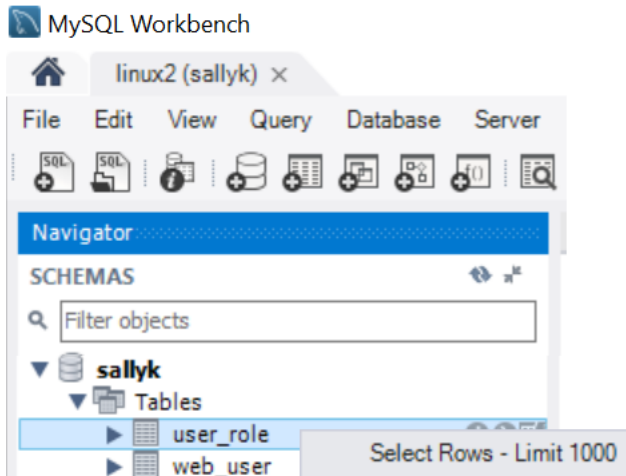
Note: if the query area is closed/disabled, right click on one of your tables and choose “Select Rows”. This will open a query window. Then replace that SQL (e.g., “SELECT * FROM web_user”) with the SQL you want to run.

- If there is an error, you should see a message at the bottom of the output window.
- Verify that you created the foreign key by right clicking on table web_user, selecting “Alter Table”, and clicking on the “Foreign Keys” tab. You should see an automatically named foreign key listed there.

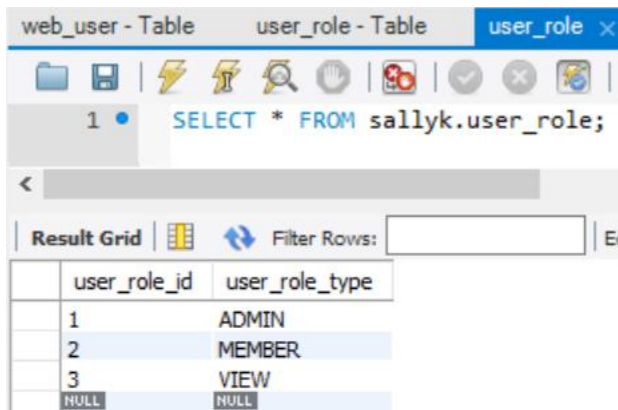


6. Add Data to Table “user_role” (and Get Primary Key Violation)

- To add data into table “user_role”, right click that table and select the top option “Select Rows”.




- Enter 2-3 records into table “user_role” then click “Apply”. Your data doesn’t have to match what’s below.



- Try to enter a record that has the same user_role_id as a record already in the table (and click “Apply”). Since this would **violate the Primary Key (PK) constraint** that you specified for the user_role.user_role_id, the database management system will give an error (like below). This is one of the main jobs of a database management system – **to prevent constraint violations, to protect the integrity of the database**.

Error: There was an error while applying the SQL script to the database.



7. Add Data to Table “web_user” (and Get Foreign Key Violation)

- Right click on table “web_user” and enter one record. Note:
 - Leave the “web_user_id” field empty. We designed this field as “auto-increment” which means we want the database management system to assign the next available (unique) number whenever we insert a record.
 - Enter the date in this format: YYYY-MM-DD.
 - Don’t put commas in the decimal field (membership_fee).
 - Enter a web_user.user_role_id value that matches one of the values in user_role.user_role_id.
 - Click “Apply” and confirm that the record got inserted.
- Now enter 3-5 more “web_user” records - all with empty web_user_id, some with empty birthday, some with empty membership_fee, all with valid user_role_id (values found in user_role.user_role_id). Click “Apply” and confirm that all records were inserted (if any one is rejected, all are rejected until you fix it).
- Try to enter a web_user record that has an invalid user_role_id (a value that is not found in user_role.user_role_id). This should cause a **foreign key constraint violation** and the database management system will give an error like this:

```
Message Log

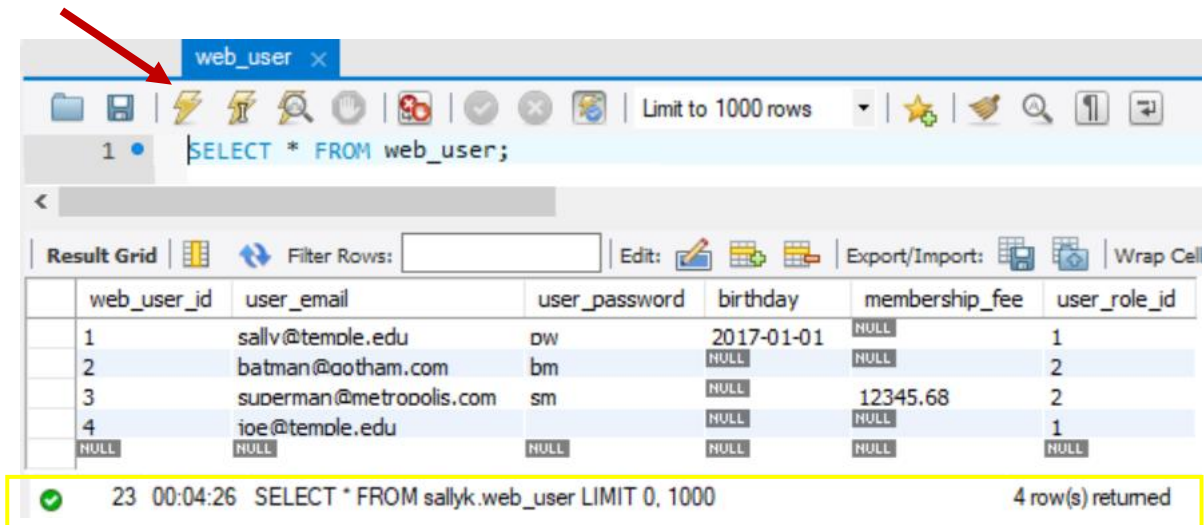
Executing:
INSERT INTO `sallyk`.`web_user` (`user_email`, `user_role_id`) VALUES ('susy', '99');

Operation failed: There was an error while applying the SQL script to the database.
ERROR 1452: 1452: Cannot add or update a child row: a foreign key constraint fails
(`sallyk`.`web_user`, CONSTRAINT `web_user_ibfk_1` FOREIGN KEY (`user_role_id`) REFERENCES
`user_role` (`user_role_id`))
SQL Statement:
INSERT INTO `sallyk`.`web_user` (`user_email`, `user_role_id`) VALUES ('susy', '99')
```

8. SELECT Data

To extract data from the database use the SQL SELECT statement. The data that is returned by running a SELECT statement is called a **result set**.

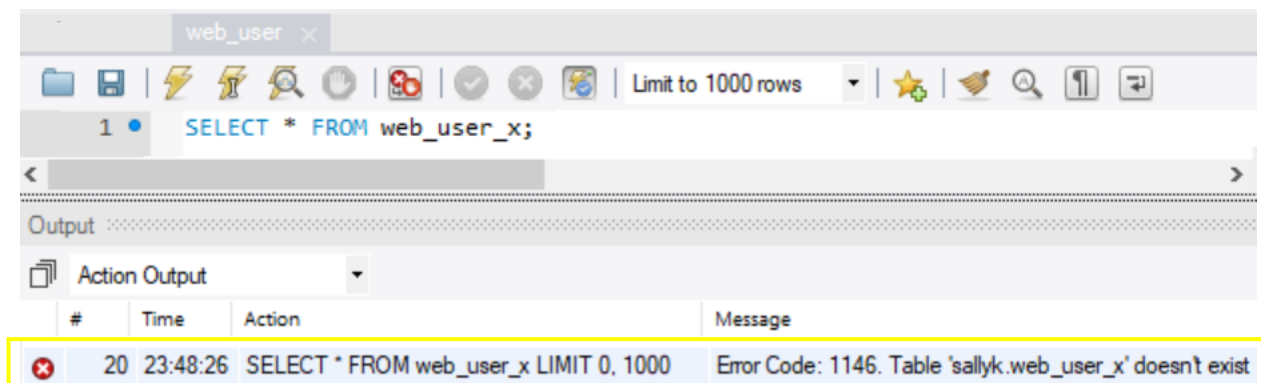
- Type your SQL into the window as shown, then click on “Query – Execute” or click on the lightning bolt icon.
 - The first statement (USE) specifies which database schema we are using. *If you have selected your schema from the “default database schema” pick list, you do not need to enter the USE statement.*
 - The second statement (SELECT) indicates that we want to see all the columns (*) from the table. Since no WHERE clause is used, all records are shown and we see a green check mark that confirms the syntactically correct SQL was executed.



The screenshot shows a database query tool interface. At the top, a tab labeled 'web_user' is active. Below the tab is a toolbar with various icons, including a lightning bolt icon for executing the query. The SQL query entered in the text area is `SELECT * FROM web_user;`. Below the query, a 'Result Grid' is displayed, showing a table with 7 columns: `web_user_id`, `user_email`, `user_password`, `birthday`, `membership_fee`, and `user_role_id`. The table contains 4 rows of data. At the bottom of the window, a status bar shows a green checkmark, the time '23 00:04:26', the query text `SELECT * FROM sallyk.web_user LIMIT 0, 1000`, and the message '4 row(s) returned'.

web_user_id	user_email	user_password	birthday	membership_fee	user_role_id
1	sallyv@temple.edu	dw	2017-01-01	NULL	1
2	batman@gotham.com	bm	NULL	NULL	2
3	superman@metropolis.com	sm	NULL	12345.68	2
4	ioe@temple.edu	NULL	NULL	NULL	1

If your SQL had a syntax error, you would see an **error message** at the bottom of the window, in the “output” area (sometimes a little hard to notice). In the example below, I misspelled the table name and you can see a very descriptive error message.



The screenshot shows the same database query tool interface, but with a syntax error. The SQL query entered is `SELECT * FROM web_user_x;`. The 'Output' area at the bottom of the window displays an error message. The error message is highlighted in a yellow box and reads: 'Error Code: 1146. Table 'sallyk.web_user_x' doesn't exist'.

#	Time	Action	Message
20	23:48:26	SELECT * FROM web_user_x LIMIT 0, 1000	Error Code: 1146. Table 'sallyk.web_user_x' doesn't exist

Instead of using *, you can specify **which columns** you want to see by listing them out.

You can specify the order you want to see in the result set by adding an **ORDER BY** clause.

The screenshot shows a database interface with a tab labeled 'web_user'. The SQL editor contains the query: `select user_email, web_user_id from web_user order by user_email;`. Below the editor, the 'Result Grid' shows the following data:

user_email	web_user_id
batman@gotham.com	2
ioe@temple.edu	4
sally@temple.edu	1
superman@metropolis.com	3

Below the result grid, the 'Output' pane shows two execution logs:

- 16 23:34:38 select web_user_id, user_email from web_user LIMIT 0, 1000 4 row(s) returned
- 17 23:36:21 select user_email, web_user_id from web_user order by user_email LIMIT 0, ... 4 row(s) returned

You can specify **which records** you want to see by adding a **WHERE** clause.

The screenshot shows the same database interface with a new SQL query: `SELECT user_email, web_user_id FROM web_user WHERE web_user_id > 2 ORDER BY user_email;`. The 'Result Grid' now shows only two rows:

user_email	web_user_id
ioe@temple.edu	4
superman@metropolis.com	3

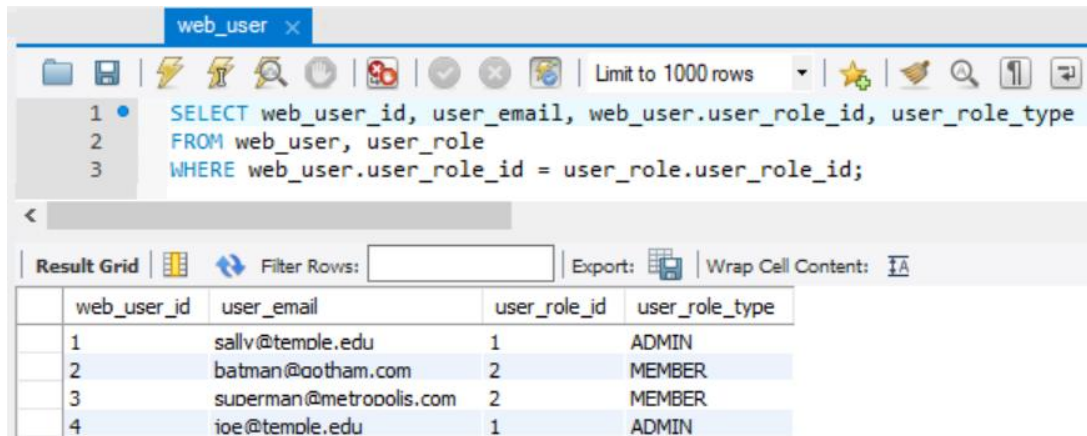
The 'Output' pane shows two execution logs:

- 17 23:36:21 select user_email, web_user_id from web_user order by user_email LIMIT 0, ... 4 row(s) returned
- 18 23:42:18 SELECT user_email, web_user_id FROM web_user WHERE web_user_id > ... 2 row(s) returned

You can split a SQL statement into multiple lines. At the end of each SQL statement, you need a “;”

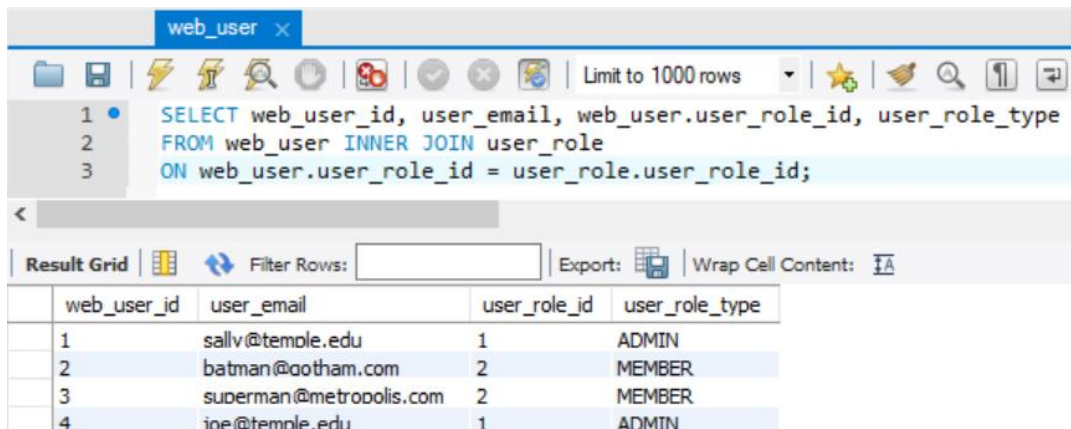
9. Join Data

Relational databases are “all about” not having any redundancy. This is why we would not put `user_role_id` and `user_role_type` directly into the `web_user` table. Instead we keep data about user_roles in the `user_role` table and simply point `web_user` records to the proper `user_role` record. To show related data, we use a JOIN. The JOIN can be written using a **WHERE clause** (as below) or an **INNER JOIN** (below that) – both yielding the same result set. When one field name is in more than one table in a SELECT statement, you must use a table name prefix for those ambiguous fields (e.g., `user_role_id` which is in both tables).



```
1 • SELECT web_user_id, user_email, web_user.user_role_id, user_role_type
2 FROM web_user, user_role
3 WHERE web_user.user_role_id = user_role.user_role_id;
```

web_user_id	user_email	user_role_id	user_role_type
1	sally@temple.edu	1	ADMIN
2	batman@gotham.com	2	MEMBER
3	superman@metropolis.com	2	MEMBER
4	ioe@temple.edu	1	ADMIN



```
1 • SELECT web_user_id, user_email, web_user.user_role_id, user_role_type
2 FROM web_user INNER JOIN user_role
3 ON web_user.user_role_id = user_role.user_role_id;
```

web_user_id	user_email	user_role_id	user_role_type
1	sally@temple.edu	1	ADMIN
2	batman@gotham.com	2	MEMBER
3	superman@metropolis.com	2	MEMBER
4	ioe@temple.edu	1	ADMIN